

CS 341: Algorithms

Module 7: Graph Algorithms

Armin Jamshidpey, Eugene Zima

Based on lecture notes by many previous CS 341 instructors

David R. Cheriton School of Computer Science, University of Waterloo

Spring 2019

Graph Algorithms

Graphs and their representation:

- A graph $G = (V, E)$ consists of a set of vertices $V = \{v_i \mid i = 1, 2, \dots, n\}$ and a set of edges $E = \{e_j \mid j = 1, 2, \dots, m\}$ that connect the vertices.
 - ▶ An edge e may be represented by the pair (u, v) where u and v are the vertices being connected by e .
 - ★ Depending on the problem, the pair may be ordered or unordered.
- Many problems can be represented as graphs:
 - ▶ Traveling Salesman Problems
 - ▶ Airline flights
 - ▶ Friends who know each other
 - ▶ Moves in a game (each node= the state of the board or game).

Directed and Undirected Graphs

- An undirected graph, $G = (V, E)$ is a pair:
 - ▶ V = set of distinct vertices.
 - ▶ E = set of edges; each member is a set of 2 vertices
 - ★ For example:

$$V = \{t, u, v, w, x, y, z\}$$

$$E = \{\{u, v\}, \{u, w\}, \{v, w\}, \{v, y\}, \{x, z\}\}$$

- A directed graph, $G = (V, E)$:
 - ▶ Is the same as an undirected graph except E is a set of ordered pair:
 - ★ For example:

$$E = \{(a, b), (a, c), (c, b), (b, e), (e, b)\} \quad (1)$$

More Definitions

- Weighted graphs:

- ▶ A weighted graph has value assigned to each of its edges.
- ▶ More formally, there is a weight function $w : E \rightarrow R$.
 - ★ $R =$ real numbers.
 - ★ Depending on the syntax being used, we may see this as $w(e)$ with $e \in E$ or $w(u, v)$ with u and v representing the vertices for the edge.

- Degree:

- ▶ The degree of vertex v , denoted by $\deg(v)$, is the number of edges that meet at v .
- ▶ Let v be a vertex in a directed graph G , The number of vertices of G adjacent **from** v is called the outdegree of v . the number of vertices of G adjacent **to** v is called the indegree of v .

Representations of Graphs

- Adjacency matrix:

- ▶ $M[i, j] = 1$ if i and j are neighbours, 0 otherwise.
- ▶ Assign each vertex an integer index (in this example, $t = 7$, $u = 2$, etc.)
- ▶ Assumes that any other info for a vertex and/or edge is in another data structure.
- ▶ The matrix is symmetric for an undirected graph.
- ▶ For a directed graph we will likely have an asymmetric matrix.

	u	v	w	x	y	z	t
u		1	1				
v	1		1		1		
w	1	1					
x						1	
y		1					
z				1			
t							

Representations of Graphs

- Adjacency matrix notes:
 - ▶ Blank row: no neighbours i.e. isolated vertex.
 - ▶ $M[i, i] = \text{self-loop}$.
 - ▶ Undirected graphs are symmetrical
- Space
 - ▶ $|V|^2$ bits
 - ▶ $(|V|^2 + |V|)/2$ (if undirected, but harder to actually implement).
 - ▶ Additional information, such as cost of an edge, could be stored in the matrix. Another option is to store a pointer to this information.
- Cost of operations
 - ▶ Are vertices i and j adjacent? $O(1)$
 - ▶ Add or delete edge: $O(1)$
 - ▶ Add vertex: increases size of matrix.
 - ▶ Find neighbours of v : $O(|V|)$.

Representations of Graphs

- Adjacency list notes:
- Space:
 - ▶ $O(|V| + |E|)$
 - ▶ Usually much smaller than $O(|V|^2)$ for a sparse graph.
- Cost of operations:
 - ▶ Add an edge $O(1)$
 - ▶ Delete an edge: Search lists for each endpoint $O(|V|)$.
 - ▶ Add vertex: Depends on the implementation.
 - ▶ Find neighbours $O(\text{number of neighbours})$ (better than Adj. Matrix)
 - ▶ Are i, j adjacent? Search the list (worse than Adj. Matrix)

More Notation

- Path:
 - ▶ Given the graph $G(V, E)$, a path from vertex u to vertex v is a sequence of vertices (v_0, v_1, \dots, v_k) such that $v = v_0$, $u = v_k$, and (v_{i-1}, v_i) is in E for all $i = 1, 2, \dots, k$.
- Simple path:
 - ▶ A path is simple if all the vertices in the path are distinct.
- Cycle:
 - ▶ A path (v_0, v_1, \dots, v_k) forms a cycle if $v_0 = v_k$, and the path contains at least one edge.
 - ▶ A graph with no cycles is acyclic.
- Connected graphs:
 - ▶ An undirected graph is connected if every pair of vertices is connected by a path.

BFS and DFS

- Breadth-first, depth-first search
- Each starts at an arbitrary node, explores the whole connected component
- Assume whole graph is connected
- General view: vertices start out coloured white (not visited)
- A visited vertex is coloured gray (visited, but may have white neighbours)
- When all neighbours of a vertex are visited, it is coloured black.

General view of searching

- Gray nodes form a “frontier”
- Can choose any neighbour of a gray node to be next visited
- In general, want to perform computation
 - ▶ Preprocess when colouring gray
 - ▶ Postprocess when colouring black
 - ▶ Analogous to tree traversal uses

Breadth-first search

- Use queue (first-in, first-out) to store gray nodes
- Start by taking any vertex, colouring it gray, add it to queue
- Repeat: find white node adjacent to head of queue, colour it gray and add it to queue
- When you can't find any such node, remove the head of the queue and colour it black

Pseudocode for BFS

```
colour_all_vertices_white()
while there is a white vertex s do
    BFS_tree(s)
```

```
BFS_tree(s: vertex)
    colour s gray (visited)
    enqueue(Q,s)
    while Q not empty do
        u <- dequeue(Q)
        for each v adjacent to u
            if v white then
                colour v gray
                enqueue(Q,v)
                 $\backslash\backslash(u,v)$  is a tree edge
            else
                 $\backslash\backslash(u,v)$  is non-tree edge
        colour u black
```

Analysis of BFS

- Assume adjacency list representation
- Each vertex enqueued once (colour changes from white to gray) and dequeued once (colour changes from gray to black)
- Body of inner while loop takes $\Theta(1)$ time
- Inner while loop implemented by scanning adjacency list of head of queue
- Therefore each edge looked at exactly twice
- Running time is $\Theta(|V| + |E|)$ or $\Theta(n + m)$

BFS trees

- BFS finds a spanning tree of the graph (edges representing first visits)
- Nontree edges are called cross edges
- A cross edge cannot connect a node to its ancestor in the tree
- A cross edge cannot connect a node to its descendant in the tree
- These are useful in proving properties of BFS searches

Single-source shortest path

- Can use BFS to compute distances $\delta(s, v)$ from source s to all other vertices v
- Compute quantity $d[v]$ in following way
 - ▶ $d[s] \leftarrow 0$
 - ▶ When adding v to queue because u is at head and there is an edge (u, v) , set $d[v] \leftarrow d[u] + 1$

Pseudocode for shortest path


```
colour_all_vertices_white()
while there is a white vertex s do
    BFS_tree(s)
```

```
BFS_tree(s: vertex)
    colour s gray (visited)
    enqueue(Q,s); d[s] <-- 0
    while Q not empty do
        u <- dequeue(Q)
        for each v adjacent to u
            if v white then
                colour v gray
                enqueue(Q,v); d[v] <-- d[u]+1
                \\\(u,v) is a tree edge
            else
                \\\(u,v) is non-tree edge
        colour u black
```


Why does this work?

Lemma 1

For any edge (u, v) , we have $\delta(s, v) \leq \delta(s, u) + 1$



Lemma 2

For all v , $d[v] \geq \delta(s, v)$.

Proof: By induction on number of enqueues

At beginning, $d[s] = 0 = \delta(s, s)$ Suppose u is being visited and adjacent white v discovered.


$$\delta(s, v) \leq \delta(s, u) + 1 \leq d[u] + 1 = d[v].$$

Why does this work?

Lemma 2 proves $d[v] \geq \delta(s, v)$.

Lemma 3

At any point, the d -values of vertices in the queue are either i or $i + 1$ for some i , and all the i values are in front of all the $i + 1$ values.




Why does this work?

Proof: By induction on number of queue operations

True at beginning (only one item in queue)

True after $u \leftarrow \text{dequeue}(Q)$

True after $\text{enqueue}(Q, v)$, because $d[v] = d[u] + 1$.



Corollary 4

d values are assigned in increasing order.

Correctness continued

Theorem 5

$$d[v] = \delta(s, v)$$

Let v be closest vertex to s with wrong d ;

$$d[v] > \delta(s, v)$$

Let u be vertex just before v on shortest $s - v$ path

$$\delta(s, v) = \delta(s, u) + 1, \quad d[u] = \delta(s, u), \quad \text{so } d[v] > d[u] + 1.$$

What colour is v when u is dequeued?

- White? Then it should have been visited from u
- Black? Then u should have been visited from v
- Gray? Then it was visited from some w , $d[v] = d[w] + 1$, and $d[w] \leq d[u]$, which implies $d[v] \leq d[u] + 1$ contradicting above inequality.

Thus no such v exists.

BFS application

- Connected components via BFS ...
- Presence of cycles ...
- Bi-partite graphs ...
- ...